

Data Systems Project

Kick-off Presentation

Linux Process Source & Sink

Connecting NebulaStream to the Linux Ecosystem

Team Members:

Saad Aldeen Mnowar • 495171

Roshni Ajay Melwani • 475317

Emi Maliqi • 498273

Aleksandar Georgiev • 468867



1. Challenge and Vision



Current State

NebulaStream primarily ingests data from static files or network sockets.

Therefore, there are limitations to direct interaction with local system tools.

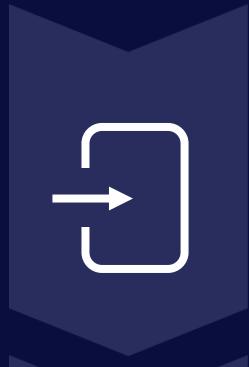


NebulaStream

Objective

Enable NebulaStream to pipe data directly to and from standard Linux processes (scripts) via Standard Streams.

2. System Architecture Overview



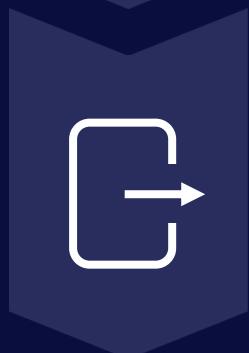
Input Side

Execute external scripts and ingest their Standard Output (stdout) directly into the streaming engine.



Processing

NebulaStream performs real-time analytics and transformations on the incoming data stream.



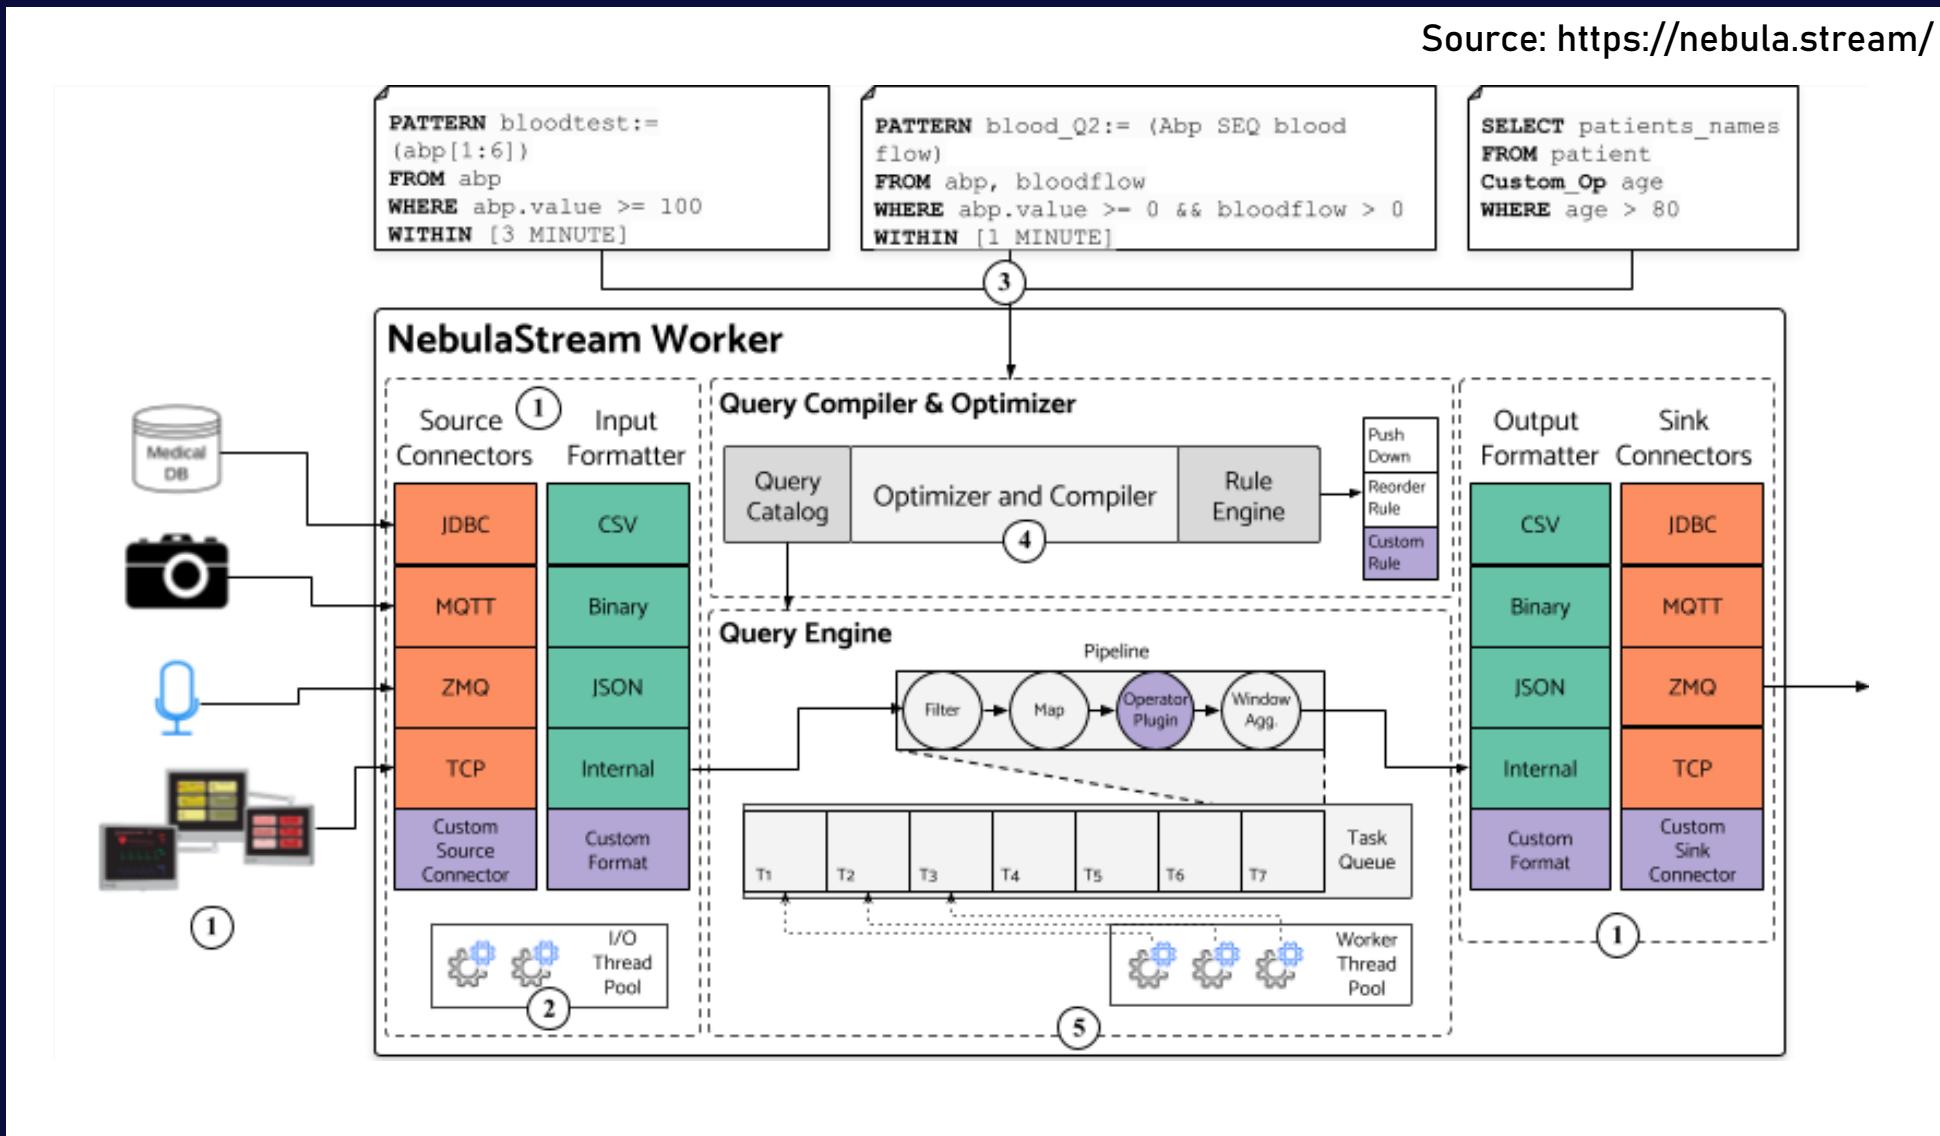
Output Side

Pipe processed results into the Standard Input (stdin) of downstream Linux tools for further use.

2. System Architecture Overview



Source: <https://nebula.stream/>



3. Implementation Plan: Source (Input) New Plugin



New Plugin Creation

Create `LinuxProcessSource` within the `nes-plugins` module as the foundation for process input handling.



Configuration Layer

Implement a robust configuration descriptor to accept arbitrary command strings (e.g., "python3 sensor.py").



Process Mechanism

Utilize `popen()` to spawn child processes and establish a read-only pipe for data capture.



Data Ingestion Logic

Implement `fillTupleBuffer` to continuously read stdout bytes and parse them into NebulaStream Tuples.

4. Implementation Plan: Sink (Output) New Plugin



New Plugin Creation

Create `LinuxProcessSink` mirroring the source structure.



Process Mechanism

Utilize `popen()` to establish a write-only pipe to the target process.



Data Formatting

Leverage existing formatters (CSV/JSON) to convert internal binary Tuples back into text streams.



Data Ingestion Logic

Implement `consumeTupleBuffer` to write formatted data directly to the child process's stdin.

5. Testing & Validation Strategy



Unit Tests

- Test individual components in isolation
- Validate buffer operations
- Verify data parsing accuracy
- Check process lifecycle management

Error Handling

- Process termination scenarios
- Pipe failure recovery
- Malformed data handling
- Resource cleanup verification

Integration Tests

- End-to-end pipeline validation
- Performance benchmarking
- Compatibility with common tools
- Stress testing under load

Thank you
for the attention!

